

Circuit safety and an Introduction to Noir

Maddiaa and Maxim
Aztec




Content

- The ZK Programming Model
- Proving System Vulnerabilities
- Common circuit bugs
- Noir
- Noir in Aztec 3
- Unconstrained Functions

We Know

- **ZK proofs are powerful**

- Information hiding

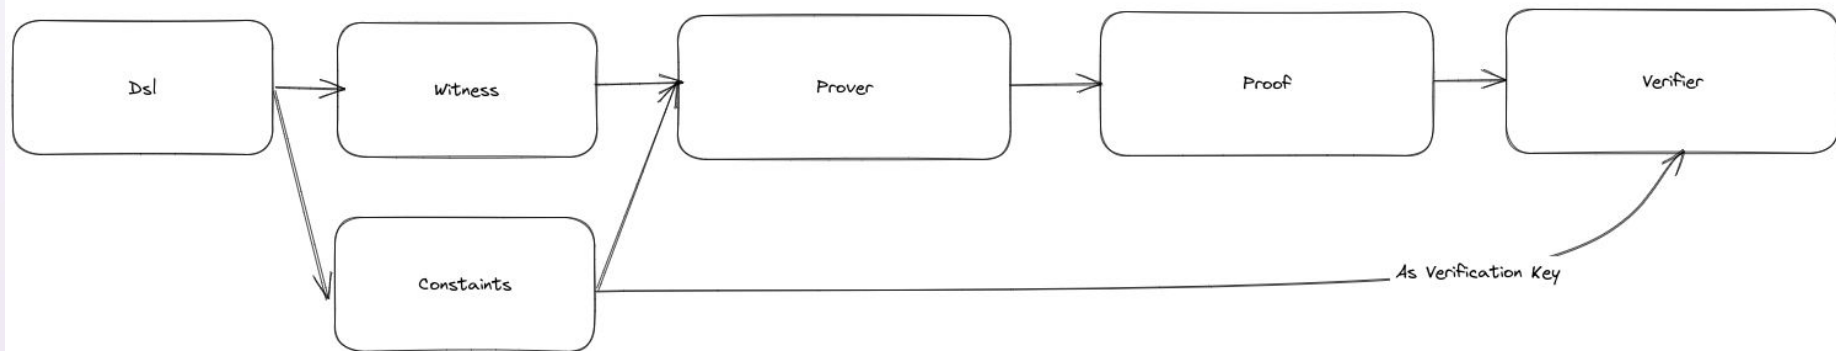


- Efficient verification



The zk programming model

- New programming paradigm
- Focus on constraining correctness (you are proving after all)



Vulnerability Classes

- Proving system vulnerabilities
- Circuit construction vulnerabilities

Proving system vulnerabilities

- Remember there is a huge dependency

Types of issues:

- Frozen Heart
- Setup Issues
- Implementation bugs

Frozen Heart

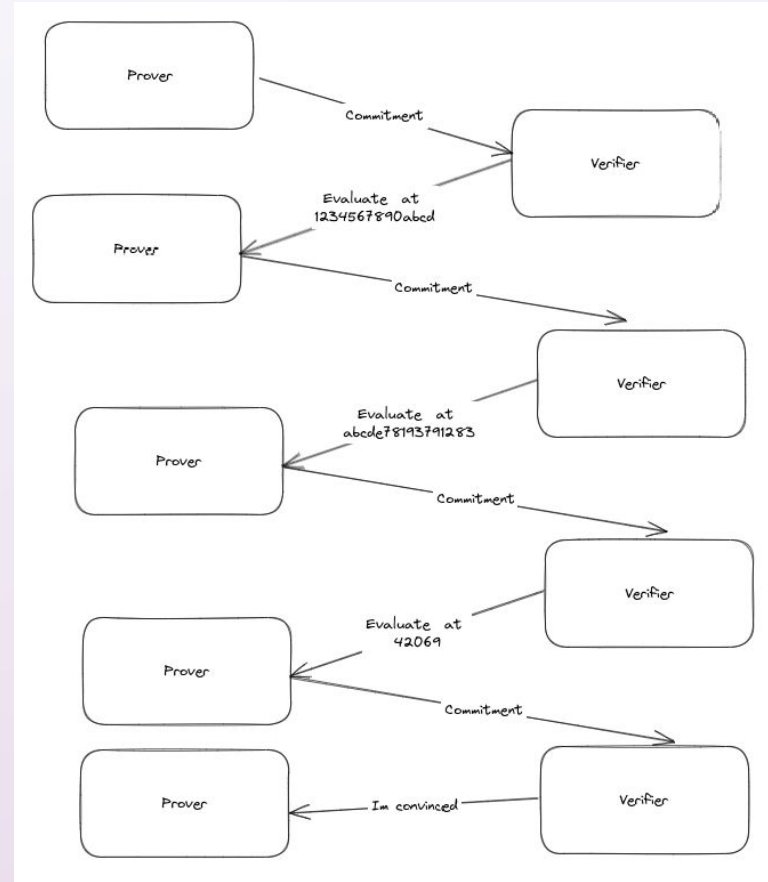
- Implementation issue in multiple plonk | bulletproof implementations
- Stems from lack of guidance over fiat shamir implementation

Frozen Heart - Fiat shamir

- Converts an interactive proof into a non interactive one.

Frozen Heart - Fiat shamir

- Interactive proofs look like this:
- The verifier will choose random values for the prover to evaluate
- Multi-round format



Frozen Heart - Fiat shamir

- A transcript is the generated proof data.

Transcript
pub input 1
pub input 2
pub input 3
w1_eval
w2_eval
w3_eval
s1_eval
s2_eval
etc.

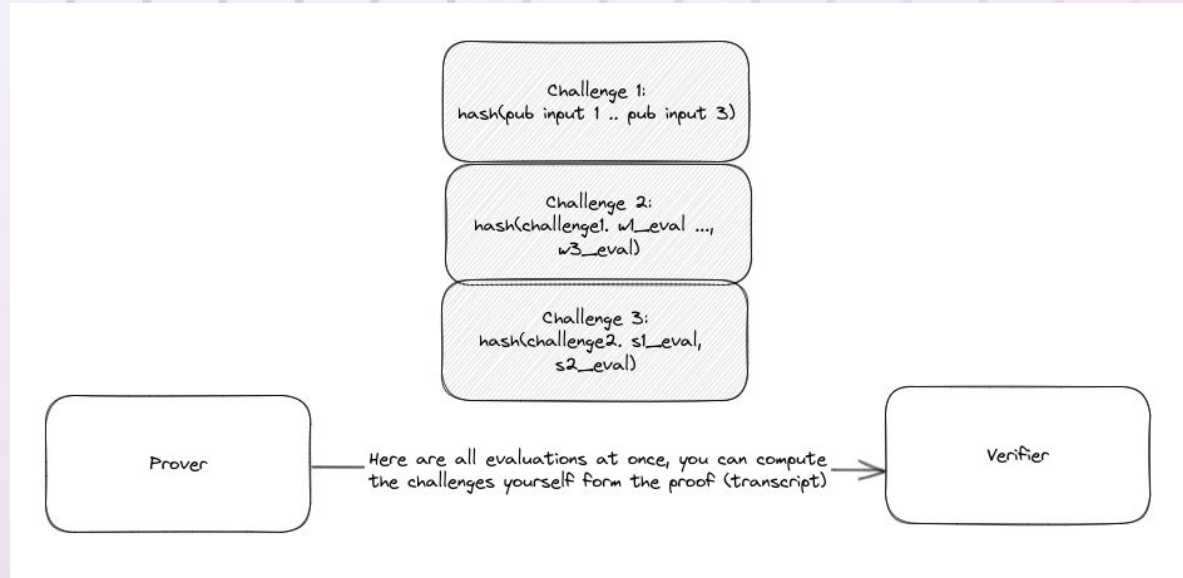
Challenge 1:
 $\text{hash}(\text{pub input 1} \dots \text{pub input 3})$

Challenge 2:
 $\text{hash}(\text{challenge1. w1_eval} \dots, \text{w3_eval})$

Challenge 3:
 $\text{hash}(\text{challenge2. s1_eval}, \text{s2_eval})$

Frozen Heart - Fiat shamir

- How to generate the challenges is known by both parties ahead of time
- Only one hop is required



Frozen Heart - The issue

- The plonk paper just assumed people would create secure constructions.
- Some implementations did not hash the entire transcript

Frozen Heart - The issue

- The plonk paper just assumed people would create secure constructions.
- Some implementations did not hash the entire transcript
- Imagine this setup

Challenge 1:
`hash(w1_eval ..., w3_eval)`

Challenge 2:
`hash(challenge1, s1_eval,
s2_eval)`

Challenge 3:
`hash(challenge2, some other
points)`

Frozen Heart - The issue

- In this construction you could technically forge proofs with any public inputs, as they are not included in the transcript.

Challenge 1:
 $\text{hash}(w1_eval \dots, w3_eval)$

Challenge 2:
 $\text{hash}(\text{challenge1}, s1_eval, s2_eval)$

Challenge 3:
 $\text{hash}(\text{challenge2}, \text{some other points})$

Frozen Heart - Overview

- Implementation issue
- It allows proof forgery, but is very hard to attack in practice. You'll still have to do a lot of hashing to create a forgery

Q&A Pause
(I'm not a cryptographer though so no hard
questions)

Common Circuit vulnerabilities

Types of issues:

- Under-constrained circuits
- Mismatching bit lengths
- Double spend nullifier issues

Smart contract circuit related issues

- Smart contract nullifier issues

Under-constrained Circuits

Underconstrained circuits - an example

- Tornado cash



Underconstrained circuits - what dis?

- The constructed circuit does not have the correct rules to fulfill its intended purpose

Underconstrained circuits - An illustration

- Me saying just add a constraint doesn't actually help much
- Why does not adding a constraint cause bugs
- A quick detour into plonk land
- (dont quote me on any of this I am no cryptogtapher)

Underconstrained circuits - An illustration

- Take a simple hash function $H(x) \Rightarrow y$
- In a plonkish circuit we will have a trace with a preimage, a trace of the hash function and an output.

Underconstrained circuits - An (very rough) illustration

- Tak
- In a
fun

	a	b	c
1 Input			x
2 Output			y
3 Hash gate 1	x	some_value_1	some_value_2
4 Hash gate 2	some_value_2	some_value_3	some_value_n-1
...			
n Hash gate n	some_value_n-1	some_value_n	y

hash

Underconstrained circuits - An (very rough) illustration

- Take a simple hash function $H(x) \Rightarrow y$
- input x
- input y
- Some hash trace that takes me from $x \rightarrow y$
- I MUST specifically constrain that $y_{in} = y_{out}$ of the hash function
- If not can create valid proofs

	a	b	c
1 Input			x
2 Output			y
3 Hash gate 1	x	some_value_1	some_value_2
4 Hash gate 2	some_value_2	some_value_3	some_value_{n-1}
...			
n Hash gate n	some_value_{n-1}	some_value_n	y

Underconstrained circuits - An (very rough) illustration

- If we do not constrain $2c == nc$, we can technically put anything in $2c$ and the verifier be convinced!

	a	b	c
1 Input			x
2 Output			y
3 Hash gate 1	x	some_value_1	some_value_2
4 Hash gate 2	some_value_2	some_value_3	some_value_n-1
...			
n Hash gate n	some_value_n-1	some_value_n	y

Copy Constraints

- $1c == 3a$
- $2c == nc$
- $3c == 2a$
- $3c == na$
- etc ...

Underconstrained circuits - an example

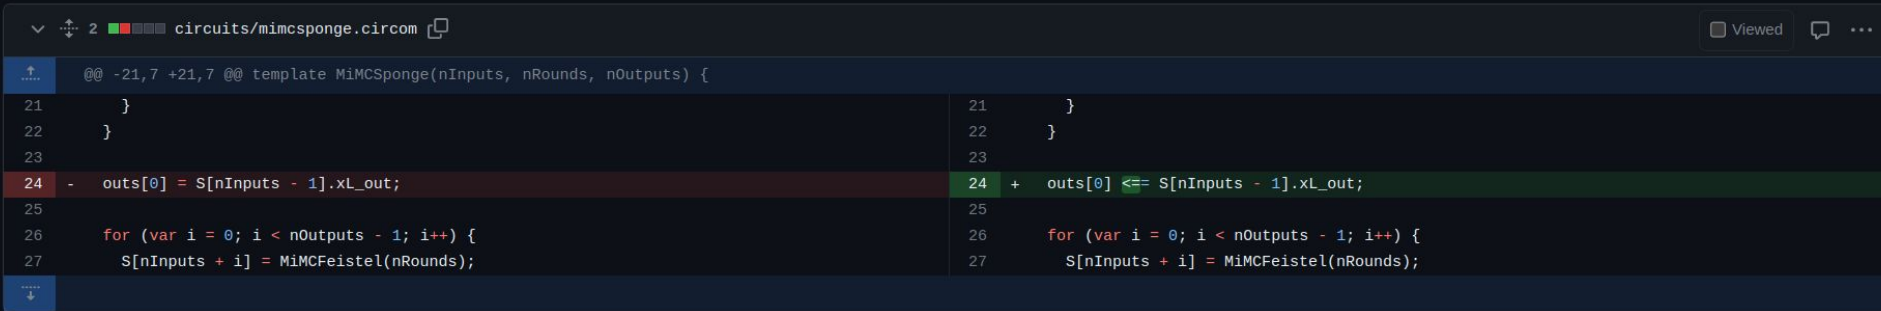
- The bug was within circomlibs standard library for a MiMC hash
- The value was assigned but not actually constrained

```
outs[0] = S[nInputs - 1].xL_out;
```

Underconstrained circuits - The fix

- Literally adding a constraint

`outs[0] <= S[nInputs - 1].xL_out;`



```
circuits/mimcsponge.circom

@@ -21,7 +21,7 @@ template MiMCSponge(nInputs, nRounds, nOutputs) {
21     }
22 }
23
24 - outs[0] = S[nInputs - 1].xL_out;
25
26 for (var i = 0; i < nOutputs - 1; i++) {
27     S[nInputs + i] = MiMCFistel(nRounds);
28 }
29 }
```

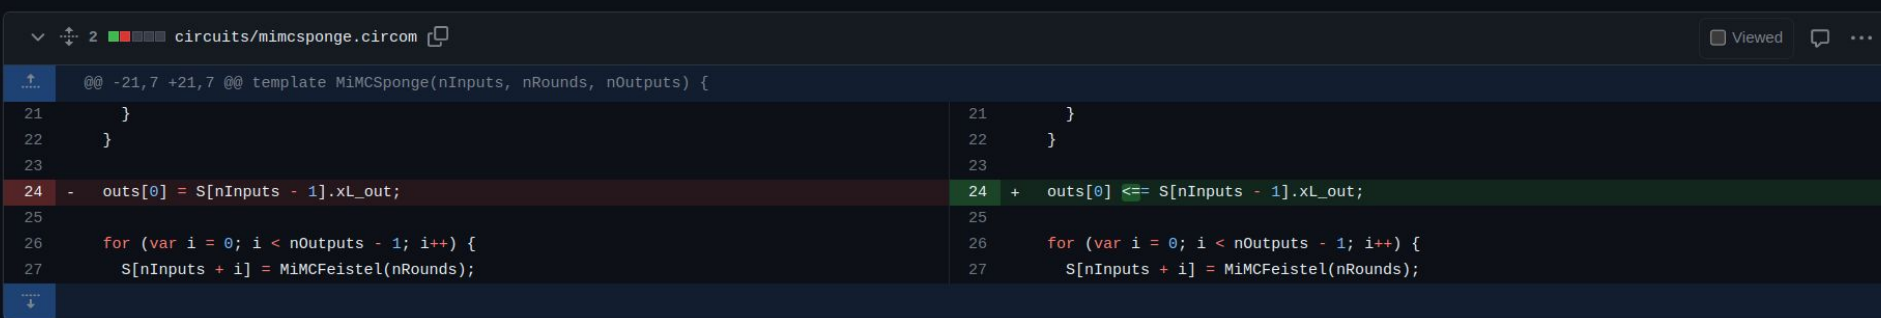
```
21     }
22 }
23
24 + outs[0] <= S[nInputs - 1].xL_out;
25
26 for (var i = 0; i < nOutputs - 1; i++) {
27     S[nInputs + i] = MiMCFistel(nRounds);
28 }
29 }
```

Underconstrained circuits - The fix

- Literally adding a constraint

P.s. noir fixes this

```
outs[0] <== S[nInputs - 1].xL_out;
```



```
circuits/mimcponge.circom

@@ -21,7 +21,7 @@ template MiMCSponge(nInputs, nRounds, nOutputs) {
21     }
22 }
23
24 - outs[0] = S[nInputs - 1].xL_out;
25
26 for (var i = 0; i < nOutputs - 1; i++) {
27     S[nInputs + i] = MiMCFistel(nRounds);
28 }
29 }
```

```
21     }
22 }
23
24 + outs[0] <== S[nInputs - 1].xL_out;
25
26 for (var i = 0; i < nOutputs - 1; i++) {
27     S[nInputs + i] = MiMCFistel(nRounds);
28 }
29 }
```

Underconstrained circuits

- Circom is a pretty low level language, you must hand wire constraints
- Static analysis tools exist
 - <https://github.com/Veridise/Picus>
 - <https://github.com/trailofbits/circomspect>
- Noir will wire these constraints for you.

The usual suspects

The usual suspects

- Overflow / underflow
- Signature malleability

Overflow / Underflow

Overflow / Underflow: An example

- EVERYTHING IS MOD SOME PRIME FIELD
- Sometimes exists in the solidity context (Make sure any circuit values used in smart contracts are mod the prime field)
- In the circuit context, add range constraints wherever you can.



Signature Malleability

DON'T USE SIGNATURES AS YOUR NULLIFIERS

- This creates non deterministic nullifiers



Nullifier: recap / primer

- State will exist in two trees
 - Data tree and Nullifier Tree





Nullifier: recap / primer

- Insert value





Nullifier: recap / primer

- Spend value
 - We don't update the data tree, we insert into the nullifier tree a value deterministically derived from some information in the data tree




Nullifier: recap / primer

- Double spend a value
 - Insert into the nullifier tree twice for a single value



Nullifier: recap / primer (bruh)

- Spend
 - We
 - from



Non Deterministic Nullifiers: example

- ECDSA construction is malleable
 - Why?
 - secp256k1 has a prime order of p
 - A signature (r, s) has another value signature $(r, p-s)$.
- This leads to a many->one rel between nullifiers (there is more than one valid nullifier for each commitment)


Non Deterministic Nullifiers: example

- This was a real world vulnerability (0xPARC stealthdrop)
- They have an excellent write up
- <https://github.com/stealthdrop/stealthdrop/blob/main/README.md#ecdsa-signature-verification>

Mismatching bit length

Vulnerability in Aztec Connect :(

- Nullifier double spend possibility
- Each nullifier had a tree 32bit index
- There was no check that the provided nullifier was 32 bits
- Only the first 32 bits were checked so a correct value appended with ANYTHING would become a valid nullifier




Current ZK language landscape

Language landscape

- Circom (low level - very customisable)

The pyramid of circuit construction



Writing circuits in cpp with aztec's proving lib (my day job help)

```
void join_split_circuit(Composer& composer, join_split_tx const& tx)
{
    join_split_inputs inputs = {
        .proof_id = witness_ct(&composer, tx.proof_id),
        .public_value = suint_ct(witness_ct(&composer, tx.public_value), NOTE_VALUE_BIT_LENGTH, "public_value"),
        .public_owner = witness_ct(&composer, tx.public_owner),
        .asset_id = suint_ct(witness_ct(&composer, tx.asset_id), ASSET_ID_BIT_LENGTH, "asset_id"),
        .num_input_notes = witness_ct(&composer, tx.num_input_notes),
        .input_note1_index = suint_ct(witness_ct(&composer, tx.input_index[0]), DATA_TREE_DEPTH, "input_index0"),
        .input_note2_index = suint_ct(witness_ct(&composer, tx.input_index[1]), DATA_TREE_DEPTH, "input_index1"),
        .input_note1 = value::witness_data(composer, tx.input_note[0]),
        .input_note2 = value::witness_data(composer, tx.input_note[1]),
        .output_note1 = value::witness_data(composer, tx.output_note[0]),
        .output_note2 = value::witness_data(composer, tx.output_note[1]),
        // Construction of partial_claim_note_witness_data includes construction of bridge_call_data, which contains
        // many constraints on the bridge_call_data's format and the bit_config's format:
        .partial_claim_note = claim::partial_claim_note_witness_data(composer, tx.partial_claim_note),
        .signing_pub_key = stdlib::create_point_witness(composer, tx.signing_pub_key),
        .signature = stdlib::schnorr::convert_signature(&composer, tx.signature),
        .merkle_root = witness_ct(&composer, tx.old_data_root),
        .input_path1 = merkle_tree::create_witness_hash_path(composer, tx.input_path[0]),
        .input_path2 = merkle_tree::create_witness_hash_path(composer, tx.input_path[1]),
        .account_note_index =
            suint_ct(witness_ct(&composer, tx.account_note_index), DATA_TREE_DEPTH, "account_note_index"),
        .account_note_path = merkle_tree::create_witness_hash_path(composer, tx.account_note_path),
        .account_private_key = witness_ct(&composer, static_cast<fr>(tx.account_private_key)),
        .alias_hash = suint_ct(witness_ct(&composer, tx.alias_hash), ALIAS_HASH_BIT_LENGTH, "alias_hash"),
        .account_required = bool_ct(witness_ct(&composer, tx.account_required)),
        .backward_link = witness_ct(&composer, tx.backward_link),
        .allow_chain = witness_ct(&composer, tx.allow_chain),
    };
    auto outputs = join_split_circuit_component(inputs);
    const field_ct defi_root = witness_ct(&composer, 0);
    defi_root.assert_is_root();
}
```


Writing circuits in circom (somehow worse than cpp imo)

```
template MerkleTreeInclusionProof(n_levels) {  
    signal input leaf;  
    signal input path_index[n_levels];  
    signal input path_elements[n_levels][1];  
    signal output root;  
  
    component hashers[n_levels];  
    component mux[n_levels];  
  
    signal levelHashes[n_levels + 1];  
    levelHashes[0] <== leaf;  
  
    for (var i = 0; i < n_levels; i++) {  
        // Should be 0 or 1  
        path_index[i] * (1 - path_index[i]) == 0;  
  
        hashers[i] = HashLeftRight();  
        mux[i] = MultiMux1(2);  
  
        mux[i].c[0][0] <== levelHashes[i];  
        mux[i].c[0][1] <== path_elements[i][0];  
  
        mux[i].c[1][0] <== path_elements[i][0];  
        mux[i].c[1][1] <== levelHashes[i];  
  
        mux[i].s <== path_index[i];  
        hashers[i].left <== mux[i].out[0];  
        hashers[i].right <== mux[i].out[1];  
    }  
}
```

Writing circuits in noir (very based)

```
use dep::std;

fn main(
  recipient: pub Field,
  // Private key of note
  // all notes have the same denomination
  priv_key: Field,
  // Merkle membership proof
  note_root: pub Field,
  index: Field,
  note_hash_path: [Field; 3],
  // Random secret to keep note_commitment private
  secret: Field,
  // Hash to be checked against the nullifier computed in the circuit
  nullifierHash: pub Field,
) -> pub Field {
  // Compute public key from private key to show ownership
  let pubkey = std::scalar_mul::fixed_base(priv_key);
  let pubkey_x = pubkey[0];
  let pubkey_y = pubkey[1];

  // Compute input note commitment
  let note_commitment = std::hash::pedersen([pubkey_x, pubkey_y, secret]);

  // Compute input note nullifier
  let nullifier = std::hash::pedersen([note_commitment[0], index, priv_key]);
  constrain nullifierHash == nullifier[0];

  // Check that the input note commitment is in the root
  let new_root = compute_root_from_leaf(note_commitment[0], index, note_hash_path);
  constrain new_root == note_root;

  // Cannot have unused variables, return the recipient as public output of the circuit
  recipient
}
```




Why Noir?

✓ Write ZK like a Dev 🧑💻
✗ not as a Mad Scientist 🧑🔬

Before Noir

- **Writing ZK programs was hard**
 - Advanced cryptography background
 - Manually write system constraints
- **Gardens were walled**
 - Tied to proving backends
 - Tied to ecosystems / chains

NOOOOO ZK DEV IS HARD ITS FOR PROS ONLY



What Noir Offers

! we're still in alpha




Spot the Difference

```
fn sum(x : u32, y : u32) -> u32 {  
  x + y  
}
```

```
fn sum(x : u32, y : u32) -> u32 {  
  x + y  
}
```


Spot the Difference



Spot the Difference

```
fn sum(x : u32, y : u32) -> u32 {  
  x + y  
}
```

Noir

```
fn sum(x : u32, y : u32) -> u32 {  
  x + y  
}
```



The Noir Programming Language

Primitive Types

- Integer
- Field
- Boolean
- String

```
fn main() {  
  let x : u32 = 10;    // Integer  
  let y : Field = 10;  // Field  
  let true = true;    // Boolean  
  let hello = "hello"; // String  
}
```

The Noir Programming Language

Compound Types

- **Arrays**
- Tuples
- Structs

```
fn main(x : Field, y : Field) {  
    let my_arr = [x, y];  
    let your_arr: [Field; 2] = [x, y];  
}
```

The Noir Programming Language

Compound Types

- Arrays
- **Tuples**
- Structs



```
fn main() {  
    let tup: (u8, u64, Field) = (255, 500, 1000);  
}
```

The Noir Programming Language

Compound Types

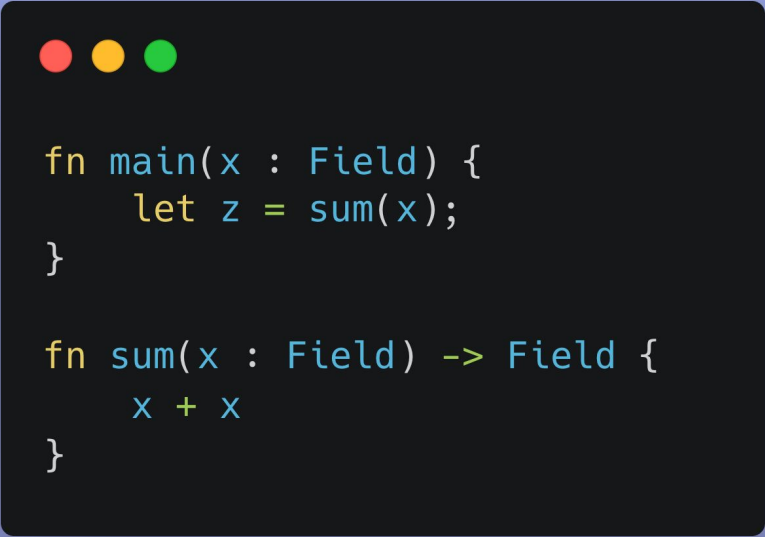
- Arrays
- Tuples
- **Structs**

```
struct Animal {  
  hands: Field,  
  legs: u8,  
  eyes: u32,  
}  
  
fn main() {  
  let dog = Animal {  
    eyes: 2,  
    hands: 0,  
    legs: 4,  
  };  
}
```

The Noir Programming Language

Modularity

- Functions



```
fn main(x : Field) {  
    let z = sum(x);  
}  
  
fn sum(x : Field) -> Field {  
    x + x  
}
```

The Noir Programming Language

Control Flow

- For loops
- If Statements

```
fn sort(mut a: [u32; 4]) -> [u32; 4] {  
    for i in 1..4 {  
        for j in 0..i {  
            if a[i] < a[j] {  
                let c = a[j];  
                a[j] = a[i];  
                a[i] = c;  
            }  
        }  
    }  
    a  
}  
  
fn main(a: [u32; 4]) {  
    let sorted_arr = sort(a);  
  
    constrain sorted_arr[0] < sorted_arr[3];  
}
```


The Noir Programming Language

Operators

- Logical
- Bitwise

```
fn main() {  
    let my_val = 5;  
    let mut shifted_val : u8 = 0;  
  
    if (my_val != 1) & (my_val < 10) {  
        shifted_val = my_val as u8 << 1;  
    }  
    constrain shifted_val == 10;  
}
```

The Noir Programming Language

And many more

- Comments
- Logging
- Generics
- First-class functions
- Dynamic array indexing
- ...

The Noir Programming Language


- **Compiles down to an IR**

- Abstract Circuit Intermediate Representation (ACIR)
- The IR can then be compiled down to any NP complete language

- **Backend Agnostic**

- Fully integrated with TurboPlonk
- Extendable to Groth16, Halo2, Gnark, etc.
- Enables proving system optimizations
 - Custom gates
 - Efficient black box functions (e.g. ECDSA, Keccak, SHA256)

The Perfect Balance



Additional Features

That the Noir tech stack offers

Tests

in Noir

```
fn add(x: u64, y: 64) -> u64 {  
    x + y  
}
```


```
#[test]  
fn test_add() {  
    constrain add(2,2) == 4;  
    constrain add(0,1) == 1;  
    constrain add(1,0) == 1;  
}
```

Q&A Pause



Noir in Aztec 3


ZK dApp Today




Problem?



ZK dApp Today



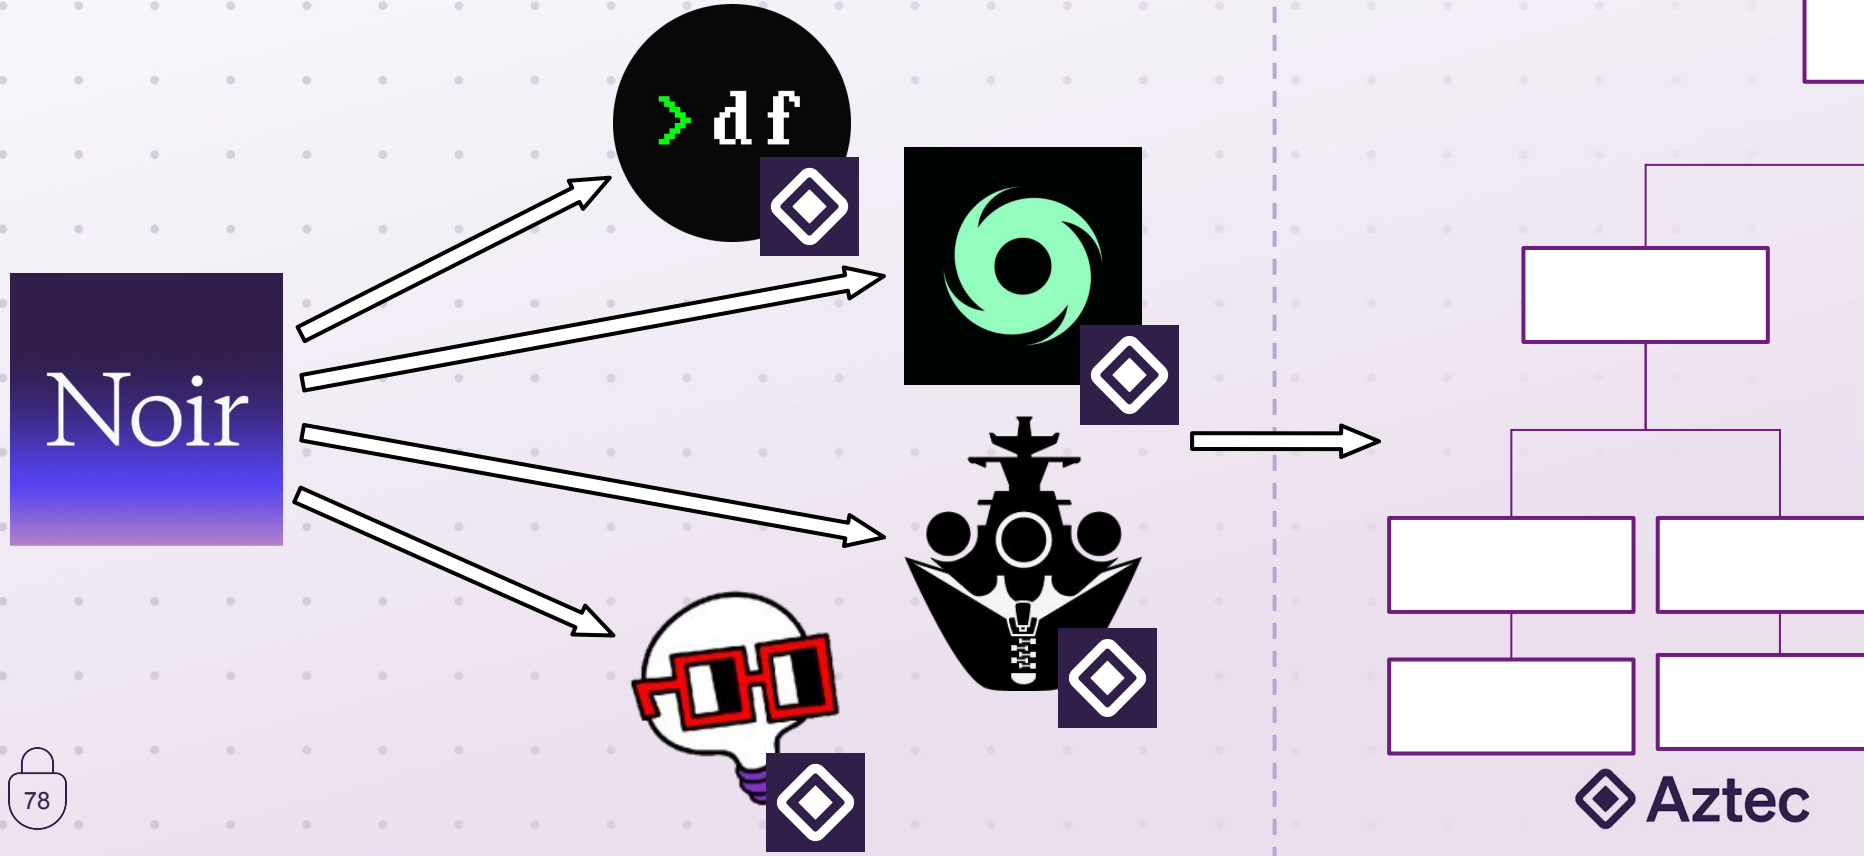
ZK dApp Today



What if...



Noir Smart Contract



Noir Smart Contract



This is soon™ Noir

This is
Aztec 3

 Aztec

Aztec 3

- **ZK-ZK Rollup**

- Inherits Ethereum's security
- Supports both private and public states
- Fully programmable
- Lower gas fees



- Noir → private + public smart contracts
- Aztec 3 → a decentralized storage and execution backend for Noir

Non-determinism

Directives

- Do not apply any constraints
- Can be thought of as an oracle
 - Fetch extra external inputs to the circuit
 - Allow one to use non-determinism
- Noir wants to prove some computation is correct
 - Not just perform the computation!
 - But it is not enough to simply give the result of the program
- In many cases, non-determinism provides a neat shortcut, but must be handled correctly

Directives Explained

- Let's decompose a Field element into an array of bytes

let x: Field = 1000;

let byte_array = x.to_le_bytes(); // 0x03E8

- Straightforward approach
 - Convert x into a u32
 - Perform bit operations to find each element of the array

```
for i in 0..FIELD_SIZE_IN_BYTES {  
    let arr[i] = (x >> (8*i)) & 0xff;  
}
```

Directives Explained

Let's decompose a Field element into an array of bytes

let x: Field = 1000;

let byte_array = x.to_le_bytes(); // 0x03E8

- Straightforward approach
 - Convert x to a u32
 - Perform bit operations to find each element of the array

**EXPENSIVE
&
INCOMPLETE**

```
for i in 0..FIELD_SIZE_IN_BYTES {  
    let arr[i] = (x >> (8*i)) & 0xff;  
}
```

WHEN SOMEONE USES
BIT OPS IN A SNARK CIRCUIT



Directives Explained

Let's use a shortcut instead!

- When generating the ACIR we will lay down the following arithmetic constraints



```
let x: Field = 1000;  
constrain [(arr[0]*2^0) + (arr[1]*2^8) + ... + (arr[31]*2^248)] - x == 0;
```

- Do we see a problem with this pseudocode?
 - How is the expression viewed by the prover?
 - How is the expression viewed by the verifier?

Directives Explained

This is non-determinism!

```
constrain (arr[0]*2^0) + (arr[1]*2^8) + ... + (arr[31]*2^248) - x == 0;
```

- The prover does not know what values of `arr` are implicitly
 - arr[0] could be 1000 and the rest of the array values are 0
 - OR the array could hold a valid byte array
- The prover must inject the correct values into the arithmetic constraint above
- Inside the directive we decompose the Field element `x` into a byte array
 - Fill in the witness values specified from ACIR generation
 - DO NOT lay down any new constraints

Unconstrained Functions

Extended Noir to Aztec 3

- Add full non-determinism to Noir
 - Users will be able to run unconstrained code
 - These functions execute outside the circuit, similarly to our existing directives
 - Users must manually constrain outputs of unconstrained circuits, using a noir program
 - We need to match values in the noir program with values from unconstrained code




Extending ACVM to Aztec 3 VM

- Unconstrained code


- Can be added within the Noir program using unsafe blocks: clean and safer!
- Can be handled by the A3 Simulator: more flexible (can do db queries, call web API, etc..)

```
secret fn transfer(amount: Field, to: Field) {  
  
    unconstrained fn reduce<NUM_NOTES: u32>(all_notes: [Note]) -> [Note; NUM_NOTES] {  
        // Imagine some code which reduces an unbounded array to a static-array  
        // ... and also ensures note.owner == msg.sender;  
    }  
  
    let my_old_notes: [Note; 2] = balances[msg.sender].get( // UTXOSet::get  
        2, // number of notes requested  
        reduce // <-- the unconstrained function defined immediately above.  
    );  
  
    // Now constrain the `reduce` logic:  
    constrain my_old_notes.all(|n| n.owner_or(msg.sender) == msg.sender);  
}
```

Links & Resources



GitHub



Docs



Awesome Noir

Reference Slides

The Noir Programming Language

- **Generics**
- First-class functions

```
struct Bar<T> {  
    one: Field,  
    two: Field,  
    other: T,  
}  
  
fn foo<T>(bar: Bar<T>) {  
    constrain bar.one == bar.two;  
}  
  
fn main(x: Field, y: Field) {  
    let bar1: Bar<Field> = Bar { one: x, two: y,  
    other: 0 };  
    let bar2 = Bar { one: x, two: y, other: [0] };  
  
    foo(bar1);  
    foo(bar2);  
}
```

The Noir Programming Language

- Generics
- **First-class functions**

```
let f = if 3 * 7 > 200 { foo } else { bar };
constrain f()[1] == 2;

// Lambdas:
constrain twice(|x| x * 2, 5) == 20;
constrain (|x, y| x + y + 1)(2, 3) == 6;

// Closures:
let a = 42;
let g = || a;
constrain g() == 42;

fn foo() -> [u32; 2] {
    [1, 3]
}

fn bar() -> [u32; 2] {
    [3, 2]
}


fn twice(f: fn(Field) -> Field, x: Field) -> Field {
    f(f(x))
}
```

Public and Private State


- Public state
 - **Account model**
 - Updateable public data tree
 - Managed by **rollup sequencers**
- Private state
 - **UTXO model**
 - Append-only private data tree
 - Append-only “indexed” nullifier tree
 - Membership proofs on **user-devices**

Noir → Smart Contracts

- Public and private storage reads/writes
- Contract scopes and silos
- Contract functions as independent ZK circuits
- Nested calls to other functions/circuits
- Calls to L1 portal contracts




Eth Nodes



Rollup Sequencers



User device



L1 functions & state changes




'Public L2' functions & state changes



'Private L2' functions & state changes

Public State

Public Data Tree




Just a value in a tree.

Accessed and updated by its key.

key: $\text{hash}(\text{contractAddress}, \text{storageSlot})$

Private State

Private Data Tree



A commitment to a value in a tree.

Inserted in next available slot.

Similar for nullifier tree...

```
hash(contractAddress, hash(storageSlot, value, owner, creator, memo, salt, inputNullifier))
```

Function Trees

- Each contract has a function tree
- A Noir *contract* function compiles to
 - ACIR Opcodes
 - Circuit
 - Verification key
- Function tree leaf
 - functionSelector
 - isPrivate
 - vkHash

Contracts Tree

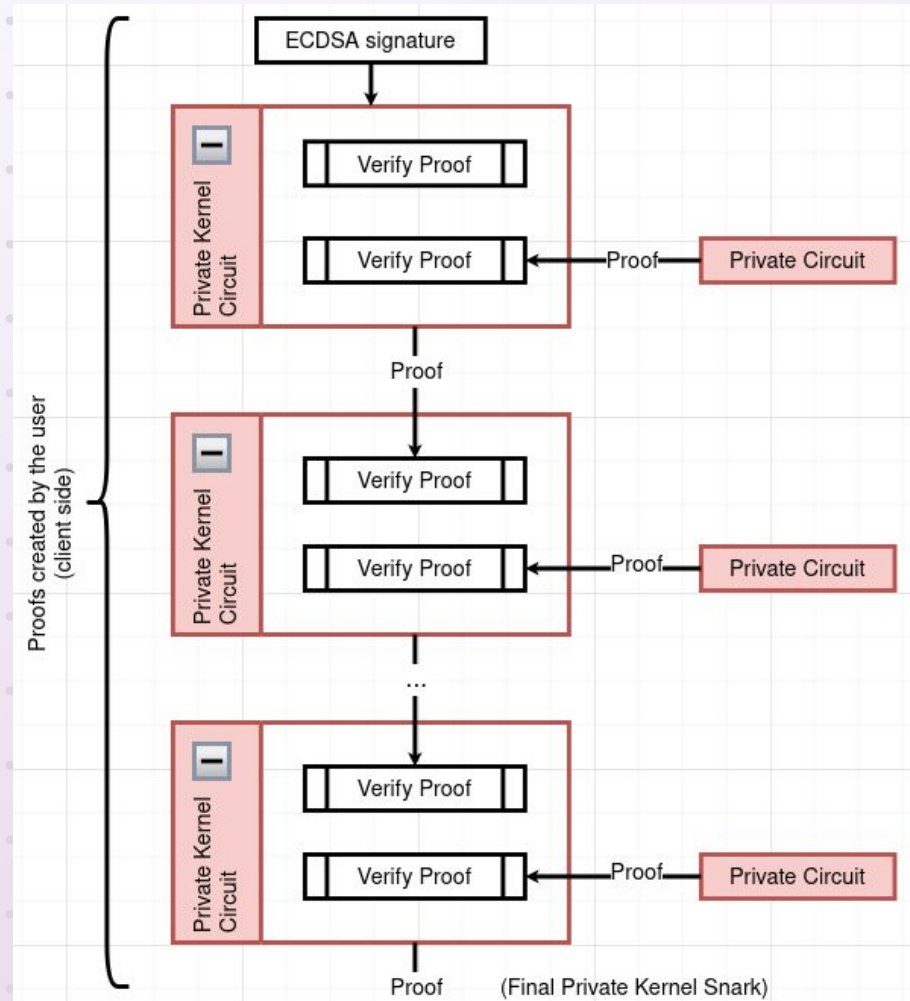
- One leaf per contract
- Append-only
- Contract tree leaf contains
 - `contractAddress`
 - `portalContractAddress`
 - `functionTreeRoot`

Kernel Circuits

- Stitch together nested function calls
- Call stack / context as public inputs
- Check function against functions tree
- Check contract against contracts tree
- Check state accesses against data trees

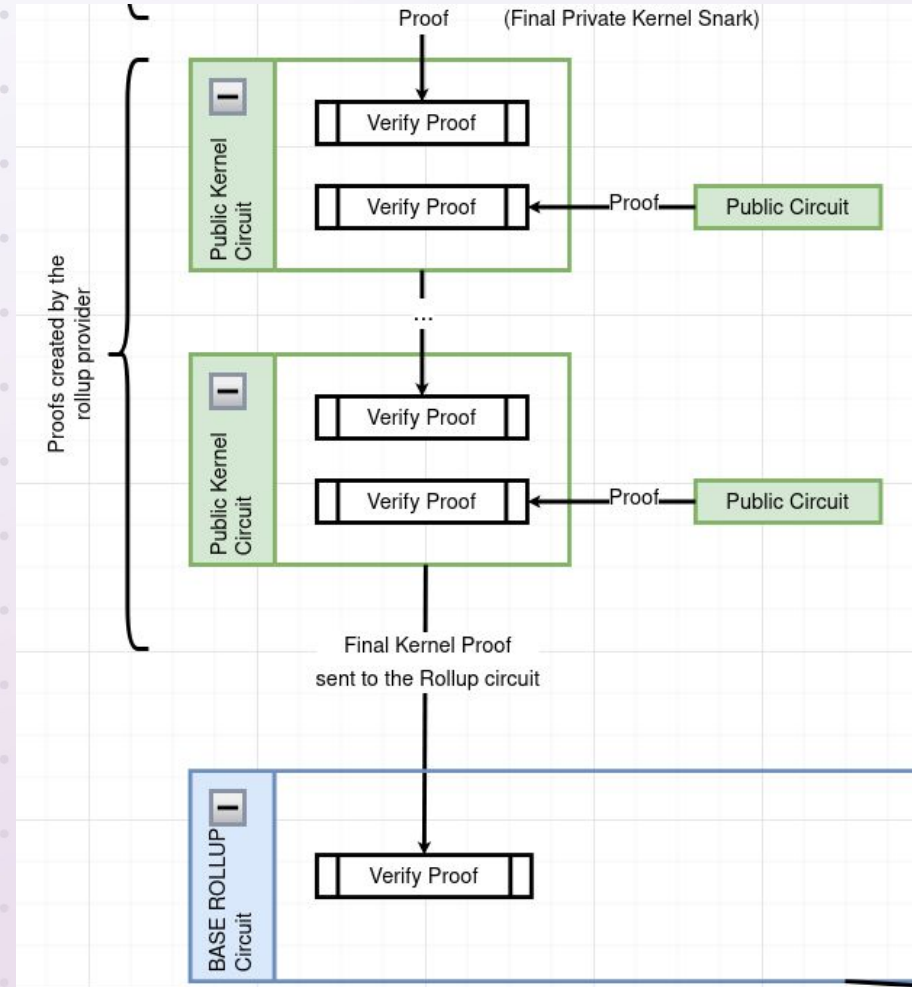
Private Kernel Circuit

- Proofs created by user
- Preserves privacy of
 - TX origin / sender
 - Function arguments
 - State accesses
- Stitch together nested private function calls
- Final proof and outputs sent to sequencer as input to public kernel



Public Kernel Circuit

- Proofs created by sequencer (actually prover)
- Accept private kernel results as input
- Preserves privacy of
 - TX origin / sender
 - All private kernel info
- Stitch together nested public function calls
- Handle L1 portal calls
- Final proof and outputs rolled up



New Noir Opcodes

- UTXO_SLOAD, UTXO_NULL, UTXO_SSTORE
- PRIV_FUNC_CALL, PUB_FUNC_CALL, L1_FUNC_CALL
- ADDRESS, PORTAL_ADDRESS
- DELEGATE_CALL, STATIC_CALL
- REVERT
- RAND
- others...